


## Local Regression

Local regression involves computing the fit at a target point  $x_0$  using only the nearby training observations.



Local regression illustrated on some simulated data, where the blue curve represents  $f(x)$  from which the data were generated, and the light orange curve corresponds to the local regression estimate  $\hat{f}(x)$ . The orange colored points are local to the target point  $x_0$ , represented by the orange vertical line. The yellow bell-shape superimposed on the plot indicates weights assigned to each point, decreasing to zero with distance from the target point. The fit  $\hat{f}(x_0)$  at  $x_0$  is obtained by fitting a weighted linear regression (orange line segment), and using the fitted value at  $x_0$  (orange solid dot) as the estimate  $\hat{f}(x_0)$ .

Local Regression at  $X = x_0$ :

1. Gather the fraction  $s = k/n$  of training points whose  $x_i$  are closest to  $x_0$ .  
*span*
2. Assign a weight  $K_{i0} = K(x_i, x_0)$  to each point in this neighborhood, so that the point furthest from  $x_0$  has weight zero, and the closest has the highest weight. All but these  $k$  nearest neighbors get weight zero.
3. Fit a weighted least squares regression of the  $y_i$  on the  $x_i$  using the aforementioned weights, by finding  $\hat{\beta}_0$  and  $\hat{\beta}_1$  that minimize


$$\sum_{i=1}^n K_{i0} (y_i - \beta_0 - \beta_1 x_i)^2$$

4. The fitted value at  $x_0$  is given by  $\hat{f}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0$ .

The span  $s$  is the tuning parameter which controls the flexibility of the non-linear fit. The smaller the value of  $s$ , the more local and wiggly will be our fit; alternatively a very large value of  $s$  will lead to a global fit to the data using all of the training observations.

We can use cross-validation to choose  $s$ , or we can specify it directly.

## Local Linear Regression



Local linear fits to the Wage data. The span specifies the fraction of the data used to compute the fit at each target point.

## Performing Splines, Smoothing Splines, and Local Regression in R


In order to fit regression splines in R, we use the `splines` library. The `bs()` function generates the entire matrix of basis functions for splines with the specified set of knots. By default, cubic splines are produced. Fitting wage to age using a regression spline is simple:

```
> library(splines)
> fit <- lm(wage ~ bs(age, knots = c(25, 40, 60)), data = Wage)
> pred <- predict(fit, newdata = list(age = age_grid), se = T)
> plot(age, wage, col = "gray")
```

```

> lines(age_grid, pred$fit, lwd = 2)
> lines(age_grid, pred$fit + 2 * pred$se, lty = "dashed")
> lines(age_grid, pred$fit - 2 * pred$se, lty = "dashed")

```



Here we have prespecified knots at ages 25, 40, and 60. This produces a spline with six basis functions. We could also use the `df` option to produce a spline with knots at uniform quantiles of the data.

```


> dim(bs(age, knots = c(25, 40, 60)))
[1] 3000    6
> dim(bs(age, df = 6))
[1] 3000    6
> attr(bs(age, df = 6), "knots")
   25%   50%   75%
33.75 42.00 51.00

```

In this case R chooses knots at ages 33.8, 42.0, and 51.0, which correspond to the 25th, 50th, and 75th percentiles of age. The function `bs()` also has a degree argument, so we can fit splines of any degree, rather than the default degree of 3 (which yields a cubic spline).

In order to instead fit a natural spline, we use the `ns()` function. Here we fit a natural spline with four degrees of freedom.

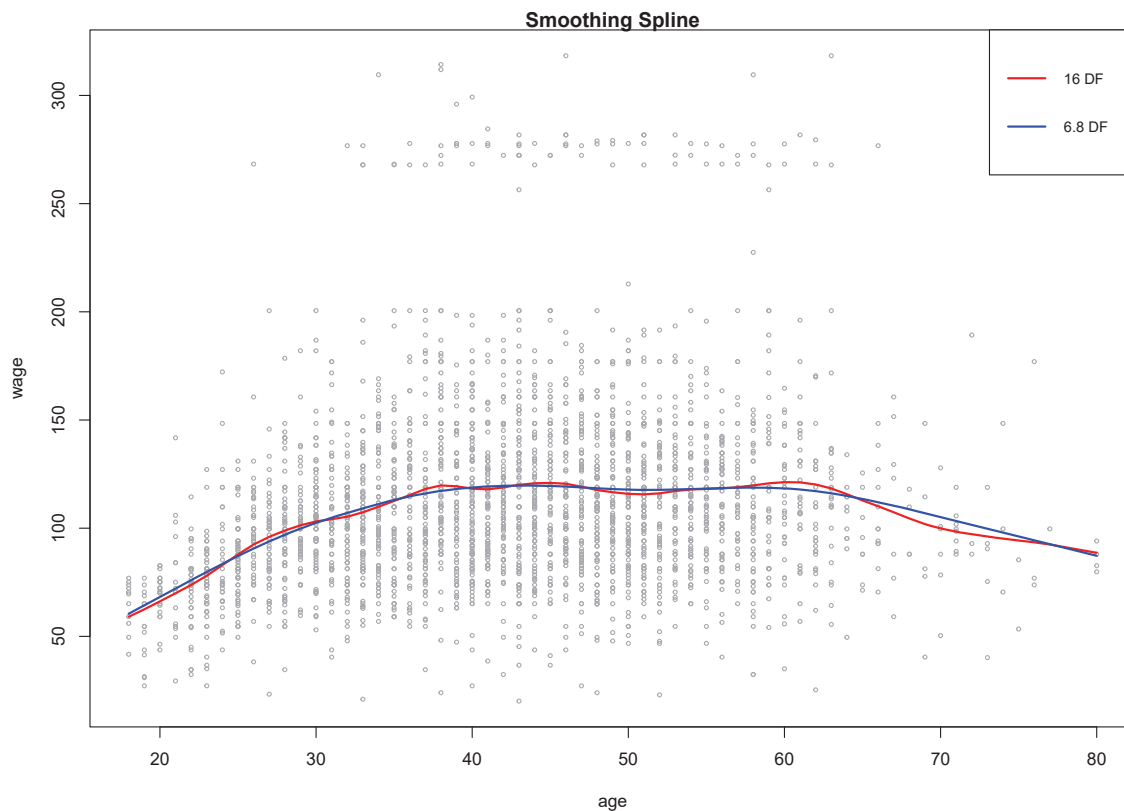
```
> fit2 <- lm(wage ~ ns(age, df = 4), data = Wage)
> pred2 <- predict(fit2, newdata = list(age = age_grid), se = T)
> lines(age_grid, pred2$fit, col = "red", lwd = 2)
```



As with the `bs()` function, we could instead specify the knots directly using the `knots` option.

In order to fit a smoothing spline, we use the `smooth.spline()` function.

```
> plot(age, wage, xlim = agelims, cex = .5, col = "darkgrey")
> title("Smoothing Spline")
> fit <- smooth.spline(age, wage, df = 16)
> fit2 <- smooth.spline(age, wage, cv = TRUE)
> fit2$df
[1] 6.794596
> lines(fit, col = "red", lwd = 2)
> lines(fit2, col = "blue", lwd = 2)
> legend("topright", legend = c("16 DF", "6.8 DF"),
      col = c("red", "blue"), lty = 1, lwd = 2, cex = .8)
```




Notice that in the first call to `smooth.spline()`, we specified `df = 16`. The function then determines which value of  $\lambda$  leads to 16 degrees of freedom. In

the second call to `smooth.spline()`, we select the smoothness level by cross-validation; this results in a value of  $\lambda$  that yields 6.8 degrees of freedom.

In order to perform local regression, we use the `loess()` function.

```
> plot(age, wage, xlim = agelims, cex = .5, col = "darkgrey")
> title("Local Regression")
> fit <- loess(wage ~ age, span = .2, data = Wage)
> fit2 <- loess(wage ~ age, span = .5, data = Wage)
> lines(age_grid, predict(fit, data.frame(age = age_grid)),
  col = "red", lwd = 2)
> lines(age_grid, predict(fit2, data.frame(age = age_grid)),
  col = "blue", lwd = 2)
> legend("topright", legend = c("Span = 0.2", "Span = 0.5"),
  col = c("red", "blue"), lty = 1, lwd = 2, cex = .8)
```



Here we have performed local linear regression using spans of 0.2 and 0.5: that is, each neighborhood consists of 20% or 50% of the observations. The larger the span, the smoother the fit.

### Generalized Additive Models

- Generalized additive models (GAMs) allow for flexible nonlinearities in several variables, but retains the additive structure of linear models.
- A natural way to extend the multiple linear regression model

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip} + \epsilon_i$$

in order to allow for non-linear relationships between each feature and the response is to replace each linear component  $\beta_j x_{ij}$  with a (smooth) non-linear function  $f_j(x_{ij})$ .

- We would then write the model as

$$y_i = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \dots + f_p(x_{ip}) + \epsilon_i$$

- We have discussed several methods for fitting functions to a single variable so far. The beauty of GAMs is that we can use these methods as building blocks for fitting an additive model.
- Take, for example, natural splines, and consider the task of fitting the model

$$\text{wage} = \beta_0 + f_1(\text{year}) + f_2(\text{age}) + f_3(\text{education}) + \epsilon$$

on the Wage data. Here year and age are quantitative variables, and education is a qualitative variable with five levels: <HS, HS, <Coll, Coll, >Coll, referring to the amount of high school or college education that an individual has completed.