


## Random Forests

- Random forests provide an improvement over bagged trees by way of a small tweak that decorrelates the trees.
- As in bagging, we build a number of decision trees on bootstrapped training samples. But when building these decision trees, each time a split in a tree is considered, a random sample of  $m$  predictors is chosen as split candidates from the full set of  $p$  predictors.
- The split is allowed to use only one of those  $m$  predictors. A fresh sample of  $m$  predictors is taken at each split, and typically we choose  $m \approx \sqrt{p}$ .
- The main difference between bagging and random forests is the choice of predictor subset size  $m$ .

If a random forest is built using  $m=p$ , then this amounts simply to bagging.

- Using a small value of  $m$  in building a random forest will typically be helpful when we have a large number of correlated predictors.
- As with bagging, random forests will not overfit if we increase  $B$ , so in practice we use a value of  $B$  sufficiently large for the error rate to have settled down.



Results from random forests for the 15-class gene expression data set with  $p = 500$  predictors. The test error is displayed as a function of the number of trees. Each colored line corresponds to a different value of  $m$ , the number of predictors available for splitting at each interior tree node. Random forests ( $m < p$ ) lead to a slight improvement over bagging ( $m = p$ ). A single classification tree has an error rate of 45.7%.

### Performing Bagging and Random Forests in R

Here we apply bagging and random forests to the Hitters data, using the `randomForest` package in R. Recall that bagging is simply a special case of a random forest with  $m = p$ . Therefore, the `randomForest()` function can be used to perform both random forests and bagging. We perform bagging as follows:

```
> library(randomForest)
> set.seed(1)
> bag_hitters <- randomForest(Log_Salary ~ .- Salary, data = Hitters,
  subset = train, mtry = 19, importance = TRUE)
> bag_hitters
```

Call:

```
randomForest(formula = Log_Salary ~ .- Salary, data = Hitters,
  mtry = 19, importance = TRUE, subset = train)
```

Type of random forest: regression

Number of trees: 500

No. of variables tried at each split: 19

Mean of squared residuals: 0.2052503

% Var explained: 73.54

The argument `mtry = 19` indicates that all 19 predictors should be considered for each split of the tree—in other words, that bagging should be done. How well does this bagged model perform on the test set?

```
> yhat_bag <- predict(bag_hitters, newdata = Hitters[-train, ])
> mean((yhat_bag - Log_Salary_test)^2)
[1] 0.190633
```

The test set MSE associated with the bagged regression tree is 0.19. We could change the number of trees grown by `randomForest()` using the `ntree` argument:

```
> bag_hitters <- randomForest(Log_Salary ~ .- Salary, data = Hitters,
  subset = train, mtry = 19, ntree = 100)
> yhat_bag <- predict(bag_hitters, newdata = Hitters[-train, ])
> mean((yhat_bag - Log_Salary_test)^2)
[1] 0.1914903
```

Growing a random forest proceeds in exactly the same way, except that we use a smaller value of the `mtry` argument. By default, `randomForest()` uses  $p/3$

variables when building a random forest of regression trees, and  $\sqrt{p}$  variables when building a random forest of classification trees. Here we use `mtry = 9`.

```
> set.seed(1)
> rf_hitters <- randomForest(Log_Salary ~ .- Salary, data = Hitters,
  subset = train, mtry = 9, importance = TRUE)
> yhat_rf <- predict(bag_hitters, newdata = Hitters[-train, ])
> mean((yhat_rf - Log_Salary_test)^2)
[1] 0.1864585
```

The test set MSE is 0.1864; this indicates that random forests yielded an improvement over bagging in this case.

Using the `importance()` function, we can view the importance of each variable.


```
> importance(rf_hitters)
```

	%IncMSE	IncNodePurity
AtBat	7.39691621	3.9659015
Hits	6.31787642	4.5657460
HmRun	0.75996667	1.3053628
Runs	5.88199191	2.3246396
RBI	6.56497811	3.0876107
Walks	8.42641065	4.5315977
Years	5.03743504	1.3763260
CAtBat	18.66803325	26.9664695
CHits	15.34037858	19.7202614
CHmRun	4.83630562	2.0156468
CRuns	10.49703146	10.5368887
CRBI	10.58221268	9.1893571
CWalks	9.70825005	6.3857389
League	-0.09445326	0.2290954
Division	2.48325541	0.2774326
PutOuts	3.02902964	1.7236172
Assists	-1.78189591	0.8043313
Errors	-1.56458058	0.7615783
NewLeague	2.20634407	0.2627862

Two measures of variable importance are reported. The first is based upon the mean decrease of accuracy in predictions on the out of bag samples. The second is a measure of the total decrease in node impurity that results from splits over that variable, averaged over all trees

Plots of these importance measures can be produced using the `varImpPlot()` function.

```
> varImpPlot(rf_hitters)
```



The results indicate that across all of the trees considered in the random forest, the number of times at bat during his career (CAtBat) and the number of hits during his career (CHits) are by far the two most important variables.

## Boosting

- Like bagging, boosting is a general approach that can be applied to many statistical learning methods for regression or classification.
- Recall that bagging involves creating multiple copies of the original training data set using the bootstrap, fitting a separate decision tree to each copy, and then combining all of the trees in order to create a single predictive model.
- Notably, each tree is built on a bootstrap data set, independent of the other trees.
- Boosting works in a similar way, except that the trees are grown sequentially: each tree is grown using information from previously grown trees.
- Boosting does not involve bootstrap sampling; instead each tree is fit on a modified version of the original data set.

### Boosting for Regression Trees:

1. Set  $\hat{f}(x) = 0$  and  $r_i = y_i$  for all  $i$  in the training set.
2. For  $b = 1, 2, \dots, B$ , repeat:
  - (a) Fit a tree  $\hat{f}^b$  with  $d$  splits ( $d+1$  terminal nodes) to the training data  $(X, r)$ .
  - (b) Update  $\hat{f}$  by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i)$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x)$$

What is the idea behind this procedure?


- Unlike fitting a single large decision tree to the data, which amounts to fitting the data hard and potentially overfitting, the boosting approach instead learns slowly.
- Given the current model, we fit a decision tree to the residuals from the model. We then add this new decision tree into the fitted function in order to update the residuals.
- Each of these trees can be rather small, with just a few terminal nodes, determined by the parameter  $d$  in the algorithm. By fitting small trees to the residuals, we slowly improve  $\hat{f}$  in areas where it does not perform well.
- The shrinkage parameter  $\lambda$  slows the process down even further, allowing more and different shaped trees to attack the residuals.

Boosting has three tuning parameters:

1. The number of trees  $B$ . Unlike bagging and random forests, boosting can overfit if  $B$  is too large, although this overfitting tends to occur slowly if at all. We use cross-validation to select  $B$ .
2. The shrinkage parameter  $\lambda$ , a small positive number. This controls the rate at which boosting learns. Very small  $\lambda$  can require using a very large value of  $B$  in order to achieve good performance. Typical values are 0.01 or 0.001, and the right choice can depend on the problem.

3. The number  $d$  of splits in each tree, which controls the complexity of the boosted ensemble. Often  $d = 1$  works well, in which case each tree is a stump, consisting of a single split. In this case, the boosted ensemble is fitting an additive model, since each term involves only a single variable. More generally  $d$  is the interaction depth, and controls the interaction order of the boosted model, since  $d$  splits can involve at most  $d$  variables.

Note: Boosting classification trees proceeds in a similar but slightly more complex way.



Results from performing boosting and random forests on the 15-class gene expression data set in order to predict cancer versus normal. The test error is displayed as a function of the number of trees. For the two boosted models,  $\lambda = 0.01$ . Depth-1 trees slightly outperform depth-2 trees, and both outperform the random forest. The test error rate for a single tree is 24%.



## Performing Boosting in R

Here we use the `gbm` package, and within it the `gbm()` function, to fit boosted regression trees to the Hitters data set. We run `gbm()` with the option `distribution = "gaussian"` since this is a regression problem; if it were a binary classification problem, we would use `distribution = "bernoulli"`. The argument `n.trees = 1000` indicates that we want 1000 trees, and the option `interaction.depth = 4` limits the depth of each tree.


```
> set.seed(1)
> Hitters_1 <- data.frame(Hitters, Log_Salary)
  #adds Log_Salary to the Hitters dataset
> boost_hitters <- gbm(Log_Salary ~ . -Salary, data = Hitters_1[train,],
  distribution = "gaussian", n.trees = 1000, interaction.depth = 4)
```

The `summary()` function produces a relative influence plot and also outputs the relative influence statistics.

```
> summary(boost_hitters)
```


	var	rel.inf
CRBI	CRBI	16.7207391
CAtBat	CAtBat	9.4574716
PutOuts	PutOuts	7.9369647
Walks	Walks	7.4876860
CRuns	CRuns	6.9584283
Assists	Assists	6.5507739
CWalks	CWalks	6.5050494
HmRun	HmRun	6.2739462
CHmRun	CHmRun	4.7390139
CHits	CHits	3.8335119
Hits	Hits	3.6066876
Runs	Runs	3.5007935
RBI	RBI	3.4535948
Errors	Errors	3.3839989
AtBat	AtBat	3.3541106
Years	Years	2.7554308
League	League	1.5925953
Division	Division	1.3084273

NewLeague NewLeague 0.5807762



We see that the number of runs batted in during his career (CRBI) is by far the most important variable. We can also produce partial dependence plots for variables. These plots illustrate the marginal effect of the selected variables on the response after integrating out the other variables.

```
plot(boost_hitters, i = "CRBI")
```



We now use the boosted model to predict log Salary on the test set

```
> yhat_boost <- predict(boost_hitters, newdata = Hitters_1[-train,],
  n.trees = 1000)
> mean((yhat_boost - Log_Salary_test)^2)
[1] 0.2447526
```

If we want to, we can perform boosting with a different value of the shrinkage parameter  $\lambda$ . The default value is 0.001, but this is easily modified.

```
> boost_hitters <- gbm(Log_Salary ~ . -Salary, data = Hitters_1[train,],
  distribution = "gaussian", n.trees = 1000, shrinkage = 0.01)
> yhat_boost <- predict(boost_hitters, newdata = Hitters_1[-train,],
  n.trees = 1000)
```

```
> mean((yhat_boost - Log_Salary_test)^2)
[1] 0.209225
```

In this case, using  $\lambda = 0.01$  leads to a lower test MSE than  $\lambda = 0.001$ .

## Bayesian Additive Regression Trees

- *Bayesian additive regression trees* (BART), is another ensemble method that uses decision trees as its building blocks.
- BART is related to bagging, random forests, and boosting approaches: each tree is constructed in a random manner as in bagging and random forests, and each tree tries to capture signal not yet accounted for by the current model, as in boosting.
- The main novelty in BART is the way in which new trees are generated.
- We let  $K$  denote the number of regression trees, and  $B$  the number of iterations for which the BART algorithm will be run.
- The notation  $\hat{f}_k^b(x)$  represents the prediction at  $x$  for the  $k$ th regression tree used in the  $b$ th iteration. At the end of each iteration, the  $K$  trees from that iteration will be summed, i.e.  $\hat{f}^b(x) = \sum_{k=1}^K \hat{f}_k^b(x)$  for  $b = 1, \dots, B$ .
- In the first iteration of the BART algorithm, all trees are initialized to have a single root node, with  $\hat{f}_k^1(x) = \frac{1}{nK} \sum_{i=1}^n y_i$ , the mean of the response values divided by the total number of trees. Thus,  $\hat{f}^1(x) = \sum_{k=1}^K \hat{f}_k^1(x) = \frac{1}{n} \sum_{i=1}^n y_i$ .
- In subsequent iterations, BART updates each of the  $K$  trees, one at a time. In the  $b$ th iteration, to update the  $k$ th tree, we subtract from each response value the predictions from all but the  $k$ th tree, in order to obtain a partial residual

$$r_i = y_i - \sum_{k' < k} \hat{f}_{k'}^b(x_i) - \sum_{k' > k} \hat{f}_{k'}^{b-1}(x_i)$$