

Week 4: Tutorial in R

Max H. Garzon



Data Life Cycle

- **Problem Definition/goal**
 - ◆ Identify/specify goals of the data analysis
 - ◆ commit to specific deliverables
- **Data pre-processing**
 - ◆ Identify appropriate data
 - ◆ Acquire data (gather, lookup, understand)
- **Data processing**
 - ◆ Identify methods (gather, cleanse, store)
 - ◆ Carry out the analysis (patterns, trends, predictions?)
- **Data post-processing**
 - ◆ Visualize and present
 - ◆ Deploy and evaluate. Iterate, if necessary



Learning Objectives

- To identify the requirements to **install and run** the package R.
- To identify the **basic structure of an R script and how to run it**
- To identify the **basic data structures** in R and learn how to use them for DS
- To identify the **basic I/O commands to enter data** for processing
- To identify the **basic I/O commands to display/visualize** to view results



Introduction – Why R ?

- R is a statistical programming language and environment for data manipulation, calculation and graphical display.
 - ◆ many useful operators for arrays and matrices.
 - ◆ many handy **tools for interactive data analysis**.
 - ◆ great **graphical facilities** for data analysis.
 - ◆ a **programming language** with conditionals, loops, user defined functions and input and output facilities



Features of R

- R is an interpreted computer language.
 - branching and looping as well as modular programming using functions.
 - user-defined functions in R are usually written in R, calling upon a smaller set of internal primitives.
 - allows user interface to procedures written in C, C++ or FORTRAN languages
 - for efficiency
 - write additional primitives



Strength of R: What R can do ?

- data handling and manipulation:
numeric, textual and many matrix operations
- high-level data analytic and statistical functions
- simple to produce great graphics
- programming language: loops, branching, subroutines
- it is free and it has a strong user-support

Weaknesses of R

- R is **not a database**, but it can be connected to DBMSs
- R is basically **a command-line interface** but some package like Rcmdr can provide nice graphical user interfaces.
- R is an **interpreted language** which can be very slow, but you can call own C/C++ code from R.
- R **lacks many spreadsheet features**, but R can input/output data from/to Excel

Data Analysis and Presentation

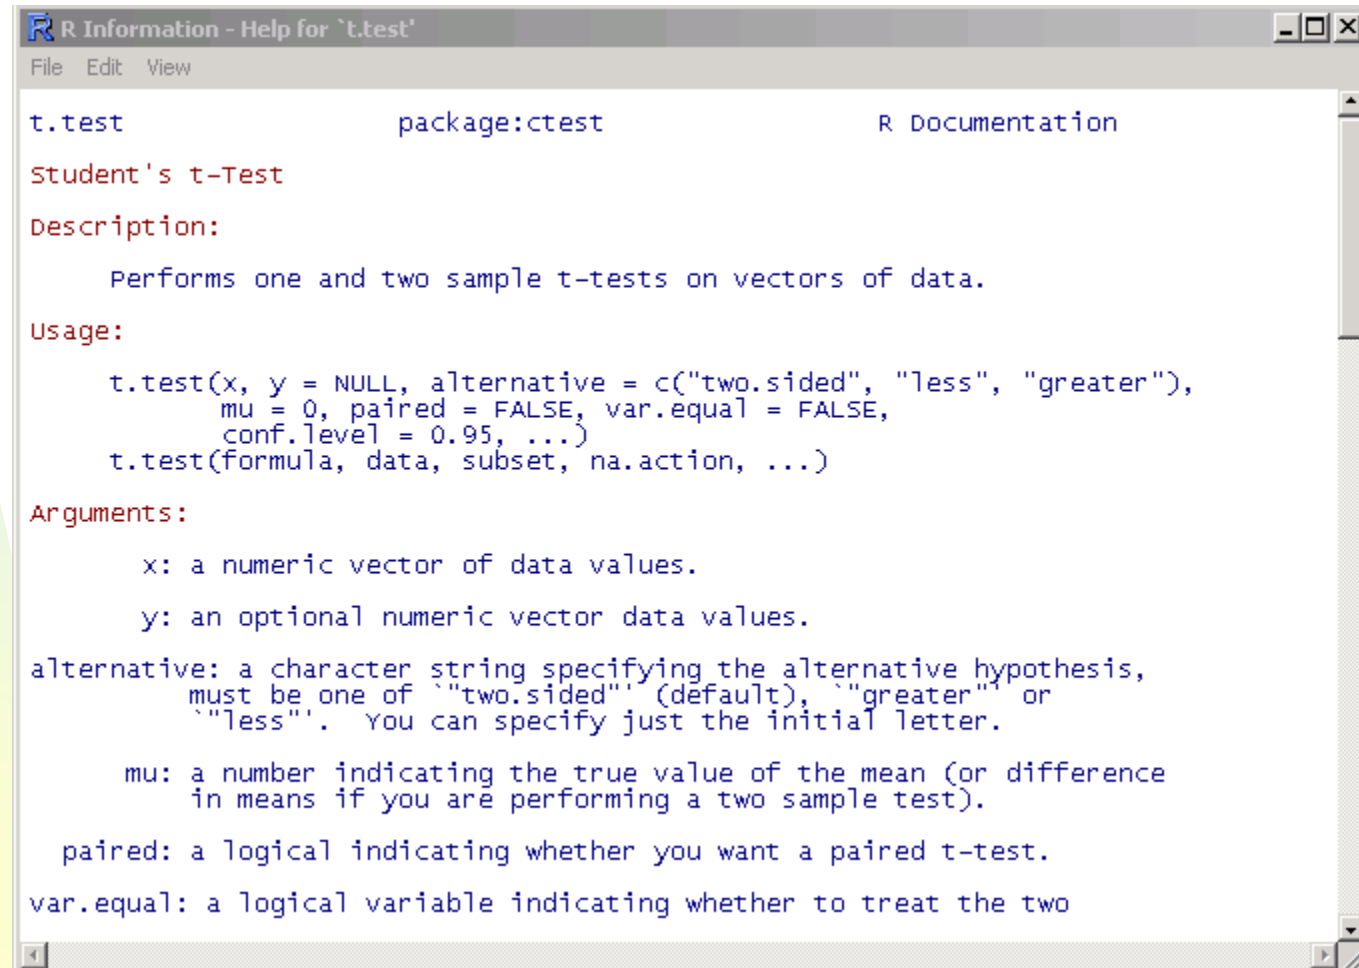
- The R distribution contains functionality for large number of statistical procedures.
 - ◆ linear and generalized linear models
 - ◆ nonlinear regression models
 - ◆ time series analysis
 - ◆ classical parametric and nonparametric tests
 - ◆ clustering
 - ◆ smoothing
- R also has a large set of functions which provide a flexible graphical environment for creating various kinds of data presentations.



Getting help

Details about a specific command whose name you know (input arguments, options, algorithm, results):

```
>? t.test  
or  
>help(t.test)
```

A screenshot of the R Help window for the 't.test' command. The window title is 'R Information - Help for `t.test`'. It shows the package 'cstest' and 'R Documentation'. The content includes the command name 't.test', a description of the Student's t-Test, and usage examples. The arguments section lists 'x', 'y', 'alternative', 'mu', 'paired', and 'var.equal' with their respective descriptions.

```
R Information - Help for `t.test`  
File Edit View  
t.test package:cstest R Documentation  
Student's t-Test  
Description:  
    Performs one and two sample t-tests on vectors of data.  
Usage:  
    t.test(x, y = NULL, alternative = c("two.sided", "less", "greater"),  
           mu = 0, paired = FALSE, var.equal = FALSE,  
           conf.level = 0.95, ...)  
    t.test(formula, data, subset, na.action, ...)  
Arguments:  
    x: a numeric vector of data values.  
    y: an optional numeric vector data values.  
alternative: a character string specifying the alternative hypothesis,  
             must be one of "two.sided" (default), "greater" or  
             "less". You can specify just the initial letter.  
mu: a number indicating the true value of the mean (or difference  
    in means if you are performing a two sample test).  
paired: a logical indicating whether you want a paired t-test.  
var.equal: a logical variable indicating whether to treat the two
```

Documentation and help file in R

- All the R functions have been documented in the form of help pages in an “output independent” form which can be used to create versions for HTML, LATEX, text etc.
 - ◆ The document “An Introduction to R” provides a more user-friendly starting point.
 - ◆ An “R Language Definition” manual
 - ◆ More specialized manuals on data import/export and extending R.



Standard packages in R

- Classical and modern statistical techniques have been implemented.
- There are several packages supplied with R (called “standard” packages) and many are available through internet sites (such as <http://cran.r-project.org>)
- `install.packages()`
lists packages available to install over the internet

Issuing commands in R

- Start R: click the icon of R after you have successfully installed the R.
- When R is started, it will prompt (`>`) and you can type in any R command.
- After you finished typing in a R command, just hit **Enter** key.
- After R finished executing your command, it will display a prompt (`>`) for your next command.
- **q()** – quits R, you will be asked whether to save workspace created.



The Workspace

- The **workspace** contains any user-defined objects that you might have created during an open session of R.
 - ◆ **Data frames, matrices, vectors, lists**
 - ◆ **Functions**
- Workspace is saved as a “.RData” file.
- You will want to know where your workspace is saved.



Working directory in R

- `getwd()` – displays current working directory
- `setwd("PATH")` – sets the working directory to PATH. Useful to work on different projects.
- ```
> getwd()
[1] "C:/Documents and
Settings/LYD/My Documents"
```
- ```
> setwd("C:/class/7150-2011/hw1")
```
- ```
> getwd()
[1] "C:/class/7150-2011/hw1"
```



# Storing data

- Every R object can be stored into and restored from a file with the commands “save” and “load”.
- This uses the XDR (external data representation) standard of Sun Microsystems and others, and is portable between MS-Windows, Unix, Mac.

```
> save(x, file="x.Rdata")
```

```
> load("x.Rdata")
```

# Managing objects in Workspace

- `ls()` – lists all objects currently in the workspace
- `rm()` – removes the object specified.
- `> ls()`  
`[1] "WD"`
- `> rm(WD) ##or rm("WD")`
- `> ls()`  
`character(0)`



# Command History

- You can save all the commands executed in R by saving your command history
- Click File, then click “Save History...”
- Choose directory where you want to save then click OK.
- Command history is saved in a “.RHistory” file
- `history()` lists last 25 commands
- `history(max.show=Inf)` lists all commands



# Built-in dataset in R

- R has many built-in datasets that you do not have to create by yourself.
- For example, R has dataset, called **mtcars**, from 1974 *Motor Trend* US magazine, for fuel consumption (**mpg**) and 10 aspects of automobile design and performance for 32 automobiles (1973–74 models).
- To see the list and description of the built-in datasets, type `data()`

# mtcars data listing

|                     | mpg  | cyl | displacement | hp  | drat | wt    | qsec  | vs | am | gear | carb |
|---------------------|------|-----|--------------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4           | 21.0 | 6   | 160.0        | 110 | 3.90 | 2.620 | 16.46 | 0  | 1  | 4    | 4    |
| Mazda RX4 Wag       | 21.0 | 6   | 160.0        | 110 | 3.90 | 2.875 | 17.02 | 0  | 1  | 4    | 4    |
| Datsun 710          | 22.8 | 4   | 108.0        | 93  | 3.85 | 2.320 | 18.61 | 1  | 1  | 4    | 1    |
| Hornet 4 Drive      | 21.4 | 6   | 258.0        | 110 | 3.08 | 3.215 | 19.44 | 1  | 0  | 3    | 1    |
| Hornet Sportabout   | 18.7 | 8   | 360.0        | 175 | 3.15 | 3.440 | 17.02 | 0  | 0  | 3    | 2    |
| Valiant             | 18.1 | 6   | 225.0        | 105 | 2.76 | 3.460 | 20.22 | 1  | 0  | 3    | 1    |
| Duster 360          | 14.3 | 8   | 360.0        | 245 | 3.21 | 3.570 | 15.84 | 0  | 0  | 3    | 4    |
| Merc 240D           | 24.4 | 4   | 146.7        | 62  | 3.69 | 3.190 | 20.00 | 1  | 0  | 4    | 2    |
| Merc 230            | 22.8 | 4   | 140.8        | 95  | 3.92 | 3.150 | 22.90 | 1  | 0  | 4    | 2    |
| Merc 280            | 19.2 | 6   | 167.6        | 123 | 3.92 | 3.440 | 18.30 | 1  | 0  | 4    | 4    |
| Merc 280C           | 17.8 | 6   | 167.6        | 123 | 3.92 | 3.440 | 18.90 | 1  | 0  | 4    | 4    |
| Merc 450SE          | 16.4 | 8   | 275.8        | 180 | 3.07 | 4.070 | 17.40 | 0  | 0  | 3    | 3    |
| Merc 450SL          | 17.3 | 8   | 275.8        | 180 | 3.07 | 3.730 | 17.60 | 0  | 0  | 3    | 3    |
| Merc 450SLC         | 15.2 | 8   | 275.8        | 180 | 3.07 | 3.780 | 18.00 | 0  | 0  | 3    | 3    |
| Cadillac Fleetwood  | 10.4 | 8   | 472.0        | 205 | 2.93 | 5.250 | 17.98 | 0  | 0  | 3    | 4    |
| Lincoln Continental | 10.4 | 8   | 460.0        | 215 | 3.00 | 5.424 | 17.82 | 0  | 0  | 3    | 4    |
| Chrysler Imperial   | 14.7 | 8   | 440.0        | 230 | 3.23 | 5.345 | 17.42 | 0  | 0  | 3    | 4    |
| Fiat 128            | 32.4 | 4   | 78.7         | 66  | 4.08 | 2.200 | 19.47 | 1  | 1  | 4    | 1    |
| Honda Civic         | 30.4 | 4   | 75.7         | 52  | 4.93 | 1.615 | 18.52 | 1  | 1  | 4    | 2    |
| Toyota Corolla      | 33.9 | 4   | 71.1         | 65  | 4.22 | 1.835 | 19.90 | 1  | 1  | 4    | 1    |
| Toyota Corona       | 21.5 | 4   | 120.1        | 97  | 3.70 | 2.465 | 20.01 | 1  | 0  | 3    | 1    |
| Dodge Challenger    | 15.5 | 8   | 318.0        | 150 | 2.76 | 3.520 | 16.87 | 0  | 0  | 3    | 2    |
| AMC Javelin         | 15.2 | 8   | 304.0        | 150 | 3.15 | 3.435 | 17.30 | 0  | 0  | 3    | 2    |
| Camaro Z28          | 13.3 | 8   | 350.0        | 245 | 3.73 | 3.840 | 15.41 | 0  | 0  | 3    | 4    |
| Pontiac Firebird    | 19.2 | 8   | 400.0        | 175 | 3.08 | 3.845 | 17.05 | 0  | 0  | 3    | 2    |
| Fiat X1-9           | 27.3 | 4   | 79.0         | 66  | 4.08 | 1.935 | 18.90 | 1  | 1  | 4    | 1    |
| Porsche 914-2       | 26.0 | 4   | 120.3        | 91  | 4.43 | 2.140 | 16.70 | 0  | 1  | 5    | 2    |
| Lotus Europa        | 30.4 | 4   | 95.1         | 113 | 3.77 | 1.513 | 16.90 | 1  | 1  | 5    | 2    |
| Ford Pantera L      | 15.8 | 8   | 351.0        | 264 | 4.22 | 3.170 | 14.50 | 0  | 1  | 5    | 4    |
| Ferrari Dino        | 19.7 | 6   | 145.0        | 175 | 3.62 | 2.770 | 15.50 | 0  | 1  | 5    | 6    |
| Maserati Bora       | 15.0 | 8   | 301.0        | 335 | 3.54 | 3.570 | 14.60 | 0  | 1  | 5    | 8    |
| Volvo 142E          | 21.4 | 4   | 121.0        | 109 | 4.11 | 2.780 | 18.60 | 1  | 1  | 4    | 2    |

# Partial listing of a dataset

- you can use `head(d, n)`, `tail(d, n)`, `print(d)` (or simply `d`) to display the first `n`, bottom `n` and all (if not too many) of the dataset `d`.

- `> head(mtcars, 2)`

|               | mpg | cyl | disp | hp  | drat | wt    | qsec  | vs | am | gear | carb |
|---------------|-----|-----|------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4     | 21  | 6   | 160  | 110 | 3.9  | 2.620 | 16.46 | 0  | 1  | 4    | 4    |
| Mazda RX4 Wag | 21  | 6   | 160  | 110 | 3.9  | 2.875 | 17.02 | 0  | 1  | 4    | 4    |

- `> tail(mtcars, 2)`

|               | mpg  | cyl | disp | hp  | drat | wt   | qsec | vs | am | gear | carb |
|---------------|------|-----|------|-----|------|------|------|----|----|------|------|
| Maserati Bora | 15.0 | 8   | 301  | 335 | 3.54 | 3.57 | 14.6 | 0  | 1  | 5    | 8    |
| Volvo 142E    | 21.4 | 4   | 121  | 109 | 4.11 | 2.78 | 18.6 | 1  | 1  | 4    | 2    |

# Special characters in R

- # #user's comment
- <- #assignment statement (also allowed:  
=          ->          <<-          ->>          )  
◆ we will use **only** <- for assignment.
- [ ] # indexing of arrays, matrices, dataframes, lists
- ( ) # encloses function input variables/arguments
- { } # groups statements (e.g. loops, functions, defs)
- ; # separates several statements on a single line
- \$ # extracting elements from lists or data frames  
“\$” is similar to “.” in other languages like C/C++/Java.



# Variable names

- Like many modern languages (C, C++, Java), the variable names are **case-sensitive**.
- While R does **not** have a concept of “**reserved words**”, several variable/function names are better treated as of “reserved words” mainly for the purpose of readability.
  - ◆ e.g. one-letter “reserved words”: `c`, `q`, `t`, `C`, `D`, `F`, `I`, and `T`.
  - ◆ `c` (concatenate), `q`(quit) , `t`(transpose of matrix), `F`(false), `T`(true), `D`(derivative), ...



# Basic data types in R

Primitive (or: atomic) data types in R are:

- numeric (integer, double, complex)
- character
- logical
- function

We can build vectors, arrays, lists from basic data types.

The primary data type in R is *vector*.

# Operators in R

- `> x <- 2 ; y <-3`
- `> x + y`  
`[1] 5`
- `> x * y`  
`[1] 6`
- `> x / y` #default is floating point division  
`[1] 0.6666667`
- `> x %/% y` # integer division  
`[1] 0`
- `> y %/% x` # integer division  
`[1] 1`
- `> x ^ y`  
`[1] 8`





# Useful functions on strings

- `paste()`
  - ◆ # concatenates and converts to string
- `substr()`, `strsplit()`
  - ◆ # substrings and splitting strings
- `grep()`, `gsub()`
  - ◆ # finds matches, replaces matches in a string
- `tolower()`, `toupper()`
  - ◆ # uppercase, lowercase conversion
- `nchar()`
  - ◆ # number of characters in string



# Example of string functions in R

- `> substr("abcdef", 2, 4)`  
`[1] "bcd"`
- `> x <- "This is a"`
- `> y <- "test only"`
- `> z <- paste(x, y); z`  
`[1] "This is a test only"`
- `> toupper(z)`  
`[1] "THIS IS A TEST ONLY"`
- `> nchar(z)`  
`[1] 19`
- `> w <- paste(z, "your score is", 90); w`  
`[1] "This is a test only your score is 90"`



# Concatenation and selection

- `> x <- c(2, 3, 4)`
- `> y <- c(6, 9, 2)`
- `> z <- c(x, y); z`  
[1] 2 3 4 6 9 2
- `> x[c(1, 3)]`  
[1] 2 4
- `> x[-2]`  
[1] 2 4

```
> length(z)
[1] 6
> x + y
[1] 8 12 6
> x / y
[1] 0.3333333 0.3333333
2.0000000
> x %/% y
[1] 0 0 2
```

# Simple functions in R

- `> x <- c(2,3,4)`
- `> sin(x)`  
[1] 0.9092974 0.1411200 -0.7568025
- `> cos(x)`  
[1] -0.4161468 -0.9899925 -0.6536436
- `> sin(x)^2+cos(x)^2 #why ? all = 1`  
[1] 1 1 1
- `> log(x)`  
[1] 0.6931472 1.0986123 1.3862944
- `> exp(x)`  
[1] 7.389056 20.085537 54.598150
- `> log10(x)`  
[1] 0.3010300 0.4771213 0.6020600



# Missing values and NaNs

R has some special values

- NA represents a missing value in the dataset
- NaN (**not a number**) because of the mathematical operations such as  $0/0$ .
- Inf (**positive infinity**) e.g.  $1/0$
- -Inf (**negative infinity**) e.g.  $\log(0)$
- NULL is an empty vector or array.

We can check them by

- `is.infinite(x)`
- `is.nan(x)`
- `is.na(x)`



# Sequence generation in R

- Common ways to generate a sequence:

- ◆ `from:to # increment  $\pm 1$ .`
- ◆ `seq(from, to, by= gap)`  
increment or length can be specified
- ◆ `rep(d, n) # replicate d n times.`

- ```
> x <- 9:5; x  
[1] 9 8 7 6 5
```
- ```
> y <- seq(0.9, 0.5, -0.1); y
[1] 0.9 0.8 0.7 0.6 0.5
```
- ```
> z <- rep(x, 2); z  
[1] 9 8 7 6 5 9 8 7 6 5
```



Logical comparisons in R

- Comparing x and y (vector or scalar) with logical comparison, it will yield a vector of True/False.
 - $x < y$, $x \leq y$
 - $\#x$ is less than, less or equal to, y
 - $x > y$, $x \geq y$
 - $\#x$ is greater than, greater or equal to, y
 - $x == y$, $x != y$
 - $\# x$ equal, not equal to, y



Logical operations in R

- We can use some logical operators for conditional expression:
 - `!, &, |, xor(x, y)`
 - # not, and, or, exclusive or
 - `any()`
 - # true if any of a vector is true
 - `all()`
 - # true if all values of a vector are true

.. Logical operations in R

- `> x <- c(1, 5, 7, 6); y <- c(2, 6, 4, 3)`
- `> x > 3 & x < 7`
[1] FALSE TRUE FALSE TRUE
- `> x <= y`
[1] TRUE TRUE FALSE FALSE
- `> x[x <= y]`
[1] 1 5
- `> (x > 3) | (y < 4)`
[1] TRUE TRUE TRUE TRUE
- `> (x > 3)`
[1] FALSE TRUE TRUE TRUE
- `> (y < 4)`
[1] TRUE FALSE FALSE TRUE
- `> (x > 3) & (y < 4)`
[1] FALSE FALSE FALSE TRUE

Vectors and arrays

- **vector** is the simplest data structure used in R which is created using `c()` function.
- **array** is an ordered collection of data of the **same** type with an **integer as its index**.
 - ◆ an array can have many dimensions.
 - ◆ matrix is simply a 2-dim array.

Using array in R

- `> x <- c(3, 5, 7, 11, 13, 19); x`
[1] 3 5 7 11 13 19
- `> y <- array(x, dim=c(2,3)); y`
[,1] [,2] [,3]
[1,] 3 7 13
[2,] 5 11 19
- `> dim(x) <- c(3,2); x`
[,1] [,2]
[1,] 3 11
[2,] 5 13
[3,] 7 19

List in R

- List in R is an object consisting of a collection of objects (components) of (possibly) different types.
- The entry of the list index is usually by some names as the key.
- It can also be referenced by its position with an integer.

Using list in R

- `> customer <- list(name="Fred", wife="Mary",
+ no.children=3, child.ages=c(4,7,9))`
- `> customer$name`
`[1] "Fred"`
- `> customer$child.ages`
`[1] 4 7 9`
- `> customer[2]`
`$wife`
`[1] "Mary"`
- `> customer[[2]]`
`[1] "Mary"`



Creating matrix

- ```
> M1 <- matrix(c(1,2,3, 11,12,13), nrow = 2,
ncol=3,
+ byrow=TRUE, dimnames = list(c("row1", "row2"),
+ c("C.1", "C.2", "C.3")))
```
- ```
> M1
```

	C.1	C.2	C.3
row1	1	2	3
row2	11	12	13
- ```
> M2 <- matrix(c(1,2,3, 11,12,13), nrow = 2,
ncol=3,
+ dimnames = list(c("row1", "row2"),
+ c("C.1", "C.2", "C.3")))
```
- ```
> M2
```

	C.1	C.2	C.3
row1	1	3	12
row2	2	11	13

Matrix operations

- ```
> M1 + M2
```

|      | C.1 | C.2 | C.3 |
|------|-----|-----|-----|
| row1 | 2   | 5   | 15  |
| row2 | 13  | 23  | 26  |
- ```
> M1 * M2 # element-wise multiplication
```

	C.1	C.2	C.3
row1	1	6	36
row2	22	132	169
- ```
> t(M2)
```

|     | row1 | row2 |
|-----|------|------|
| C.1 | 1    | 2    |
| C.2 | 3    | 11   |
| C.3 | 12   | 13   |
- ```
> M1 %*% t(M2) # multiplication
```

	row1	row2
row1	43	63
row2	203	323

Other matrix functions/op.

- `dim(A)`
 - ◆ #returns dimension of matrix or array A
- `nrow(A)` , `ncol(A)` ,
`NROW(A)` , `NCOL(A)`
 - ◆ #number of rows and columns of matrix A
- `rownames(A)` , `colnames(A)`
 - ◆ #names of rows and columns of matrix A
- `%*%`
 - ◆ # matrix multiplication



Other matrix functions/op.

- `t(A)` # transpose of matrix A
- `solve(A)` # inverse of matrix A
- `svd(A)` , `qr(A)` , `chol(A)`
 - ◆ # singular value, QR, cholesky decomposition of matrix A
- `eigen(A)` , `det(A)`
 - ◆ # eigenvalues and eigenvectors, determinant of matrix A



Combining matrices and arrays

- `cbind(x, y)`
- # binds matrices, dataframes,... columnwise
- `rbind(x, y)`
- # binds matrices, dataframes,... rowwise

```
> x <- c(1, 2, 7, 9); y <- 5:8
> cbind(x, y)
      x y
[1,] 1 5
[2,] 2 6
[3,] 7 7
[4,] 9 8
> rbind(x, y)
      [,1] [,2] [,3] [,4]
x         1     2     7     9
y         5     6     7     8
> c(x, y)
[1] 1 2 7 9 5 6 7 8
```

Data frames

data frame is a rectangular table with rows and columns; data within each column has the same type (e.g. number, text, logical), but different columns may have different types.

data.frame():

- an R command to create data frames, tightly coupled collections of variables which share many of the properties of matrices and of lists,
- used as the fundamental data structure by most of R's modeling software.

Creating data frames

- You can recreate a data frame from scratch by

```
my_data <- edit(data.frame())
```

that you can enter data into the given form.

- You can also import from external file (to be discussed later) or you can save the data created.

- ```
> my_data <- data.frame(x,y); my_data
```

|   | x | y |
|---|---|---|
| 1 | 1 | 5 |
| 2 | 2 | 6 |
| 3 | 7 | 7 |
| 4 | 9 | 8 |



# Data subsetting in R

- `x[n] , x[-n]` # select nth element, all but nth element from vector x
- `x[1:n] , x[-(1:n)]` # select first n elements, all but first n elements from x
- `x[c(1, 4, 6)]` # select element 1,4 and 6 from vector x
- `x[x>3 & x<5]` # select elements that meet condition
- `which(x==3)` # returns indices to values x that meet the condition



# .. Data subsetting in R

- `> x <- c(2, 5, 7, 11, 13, 17)`
- `> x[3]`  
`[1] 7`
- `> x[-3]`  
`[1] 2 5 11 13 17`
- `> x[1:3]`  
`[1] 2 5 7`
- `> x[-(1:3)]`  
`[1] 11 13 17`
- `> x[c(1,4,6)]`  
`[1] 2 11 17`
- `> x[-c(1,4,6)]`  
`[1] 5 7 13`
- `> which(x==13)`  
`[1] 5`



# Subsetting matrix/data frame in R

- Same rule for vector subsetting can be used for matrix or data frame (to be discussed later)
- `A[i, j], A[, j], A[i, ]` # selects element i,j, the jth column, i-th row from matrix A
- `A[, cols]` # selects columns **cols** from matrix A
- `A["name", ]` # selects row named **"name"** from matrix A
- `D$name, D[["name"]]` # selects column named **"name"** from data frame D



# ..Subsetting matrix/data frame in R

- `> mtcars[1:4,1:5]`

|        |         | mpg  | cyl | disp | hp  | drat |
|--------|---------|------|-----|------|-----|------|
| Mazda  | RX4     | 21.0 | 6   | 160  | 110 | 3.90 |
| Mazda  | RX4 Wag | 21.0 | 6   | 160  | 110 | 3.90 |
| Datsun | 710     | 22.8 | 4   | 108  | 93  | 3.85 |
| Hornet | 4 Drive | 21.4 | 6   | 258  | 110 | 3.08 |

- `> mtcars[1:4,1]`

```
[1] 21.0 21.0 22.8 21.4
```

- `> mtcars[1:4, "mpg"]`

```
[1] 21.0 21.0 22.8 21.4
```

- `> mtcars$mpg[1:4]`

```
[1] 21.0 21.0 22.8 21.4
```





# if statements in R

- `if, else, else if`
  - ◆ `#conditionally execute statements`
  - ◆ `#useful only when comparing two values, not two vectors (why not?)`
  - ◆ `# often used with all() or any()`
- R example:
  - `if(all(x < 0)) cat("all x values are negative\n")`

# ifelse statement in R

- `ifelse(cond, yes, no)`
  - ◆ `# if (component-wise) condition is true/false, executes (component-wise) statement 'yes'/'no'`
- R example:
- `x <- c(6:-4)`
- `sqrt(x) #- gives warning`
- `sqrt(ifelse(x >= 0, x, NA)) # no warning`

# Repetitive execution

- `for (el in seq) {expr}`
  - ◆ `#repeat expr for each element in seq`
- `while (cond) {expr}`
  - ◆ `#repeat expression while condition is true`
  - ◆ `#be very careful for vector comparison`
- `repeat {expr}`
  - ◆ `#repeat until break encountered`

# Breaking repetitive execution

- break
  - ◆ it terminates execution of for, while, repeat loops
  - ◆ it can be used to terminate any loop, possibly abnormally.
- next
  - ◆ it transfers execution to next iteration in loops
  - ◆ it can be used to discontinue one particular cycle and skip to the “next”.

# Example

```
> s <- 0
> for(i in 1:4) {
 s <- s + i^0.5
 print(s)
}
[1] 1
[1] 2.414214
[1] 4.146264
[1] 6.146264
```

```
> i <- 1; s <- 0;
> while(s <= 10) {
 s <- s + i^0.5
 print(s); i <- i+1
}
[1] 1
[1] 2.414214
[1] 4.146264
[1] 6.146264
[1] 8.382332
[1] 10.83182
```

# User-defined functions

## Example:

```
f <- function(a, b)
{
 return (a+b)
}
```

## Note:

- Note that `return` is a function in R; its argument must be contained in parentheses.
- The use of `return` is optional; otherwise the value of the last line executed in a function is its return value.

# apply( arr, margin, fct )

Apply the function fct along some dimensions of the array arr, according to margin, and return a vector or array of the appropriate size.

```
> x
 [,1] [,2] [,3]
[1,] 5 7 0
[2,] 7 9 8
[3,] 4 6 7
[4,] 6 3 5
> apply(x, 1, sum)
[1] 12 24 17 14
> apply(x, 2, sum)
[1] 22 25 20
```

# `lapply(li, function )`

- To each element of the list `li`, the function *function* is applied.

```
> li <- list("This", "example", "is", "great")
```

```
> lapply(li, toupper)
```

```
[[1]]
```

```
[1] "THIS"
```

```
[[2]]
```

```
[1] "EXAMPLE"
```

```
[[3]]
```

```
[1] "IS"
```

```
[[4]]
```

```
[1] "GREAT"
```



## .. `sapply( li, fct )`

`sapply` is a simplified version of `lapply` by default returning a vector or matrix if appropriate

```
> li <- list("This", "example", "is", "great")
> sapply(li, toupper)
 [1] "THIS" "EXAMPLE" "IS" "GREAT"
> fct = function(x) { return(c(x, x*x, x*x*x)) }
> sapply(1:5, fct)
```

|      | [,1] | [,2] | [,3] | [,4] | [,5] |
|------|------|------|------|------|------|
| [1,] | 1    | 2    | 3    | 4    | 5    |
| [2,] | 1    | 4    | 9    | 16   | 25   |
| [3,] | 1    | 8    | 27   | 64   | 125  |

# Input output in R

- By default, the input is from the keyboard and output is to the screen. However, there are many other methods can be used.
- `write.table(x, file)`
  - ◆ # writes object x as a dataframe to a table
- `read.table(file)`
  - ◆ # reads table from space-delimited file, aligned in columns
- `read.csv(file), read.delim(file)`
  - ◆ # reads table comma- delimited or tab-delimited file



# Reading data from files

| Price | Floor | Area | Rooms | Age | Cent.heat |
|-------|-------|------|-------|-----|-----------|
| 52.00 | 111.0 | 830  | 5     | 6.2 | no        |
| 54.75 | 128.0 | 710  | 5     | 7.5 | no        |
| 57.50 | 101.0 | 1000 | 5     | 4.2 | no        |

...

- HousePrice <-  
read.table("houses.data", header=TRUE)

# Importing and exporting data

There are many ways to get data into R and out of R. Most programs (e.g. Excel), as well as humans, know how to deal with rectangular tables in the form of tab-delimited text files.

```
> x = read.delim("filename.txt")
```

```
also: read.table, read.csv
```

```
> write.table(x, file="x.txt", sep="\t")
```

# Importing data

- Type conversions: by default, the read functions try to guess and autoconvert the data types of the different columns (e.g. number, factor, character).
  - There are options `as.is` and `colClasses` to control this – *read the online help*
  - Understand the conventions your input files use and set the quote options accordingly.

# Further Topics

Some of the topics listed will be discussed in the later modules.

- Graphics in R (page 25 in [MB], much more on Chapter 15)
- Lattice graphics ( page 30 in [MB], skip)
- Finer graphic parameter settings (page 27, [MB])
- Customized options setting (page 34, [MB])



# Questions?

